



Enterprise AI Strategy & Alignment Briefing

Office of the Chief Technology Officer | Fiscal Year 2026 | Prepared for the Executive Leadership Council

Executive Summary

Over the past eighteen months, Meridian Apex Industries has moved from exploratory pilots to a coordinated, enterprise-wide artificial intelligence program. What began in early 2025 as a handful of isolated proofs of concept inside the Engineering and Customer Support organizations has matured into a governed portfolio of twenty-three active initiatives, a dedicated AI Platform Group of seventy-one full-time engineers, and an executive-sponsored Responsible AI Council that reviews every production deployment. This briefing summarizes where we stand, how our investments are tracking against the board-approved strategy, and the alignment principles that will guide the next phase of scaling.

Our central thesis remains unchanged: AI is not a standalone product line for Meridian Apex, but a horizontal capability that compounds the value of every existing business unit. Accordingly, we measure success not by the number of models we train, but by cycle-time reduction, quality improvement, and risk reduction in the workflows our people already perform. In FY2025 the program delivered an estimated \$11.2M in annualized productivity gains against \$20.3M of cumulative investment, and we project the portfolio to reach break-even on a run-rate basis in Q4 2026. The investment curve below reflects a deliberate acceleration approved by the Finance Committee in November 2025.




Figure 1. Quarterly AI investment (bars, left axis) and dedicated AI headcount (line, right axis), Q1 2025 – Q2 2026. Source: FP&A; consolidated actuals, May 2026.

Strategic Alignment Pillars

1. Workflow-first deployment. Every initiative must attach to a named business workflow with a baseline metric. We explicitly deprioritize 'technology looking for a problem' efforts; proposals without a measurable workflow owner are returned by the intake committee. **2. Human accountability.** AI systems at Meridian Apex draft, recommend, and accelerate — they do not autonomously approve, deny, or transact where customer or employee outcomes are at stake. A named human owner signs off on every consequential decision path, and this principle is contractually embedded in our vendor agreements.

3. Platform over point solutions. The AI Platform Group operates a shared gateway for model access, retrieval infrastructure, evaluation harnesses, and audit logging. Business units consume these as internal services rather than procuring shadow tooling, which keeps our security posture consistent and our spending negotiable at enterprise scale. **4. Responsible AI by design.** Bias evaluation, data-provenance review, and red-team testing are stage gates in the delivery lifecycle, not after-the-fact audits. The Responsible AI Council, chaired by the General Counsel, holds veto authority over production launches.

Governance & Responsible AI Posture

The Responsible AI Council met eleven times during the trailing twelve months and reviewed nineteen launch requests, approving fourteen, returning four for additional evaluation evidence, and rejecting one. The most common deficiency remains incomplete documentation of training-data provenance for third-party models, and we have responded by standardizing a model-intake questionnaire that vendors must complete before procurement proceeds. In parallel, the Internal Audit function has added AI-system controls to its annual plan, with the first horizontal audit of the model gateway scheduled for August 2026. Employee enablement continues to scale: 4,800 employees (62% of the workforce) have completed the mandatory 'Working Responsibly with AI' curriculum, and 410 practitioners hold the advanced builder certification required to ship AI features to production.

Initiative Portfolio — Tier-1 Programs

Initiative	Business Owner	Stage	FY26 Budget	Annualized Benefit	Status
Atlas — engineering copilot rollout	VP, Engineering	Scaled	\$2.4M	\$4.1M	On track
Beacon — support deflection & triage	SVP, Customer Ops	Scaled	\$1.8M	\$3.3M	On track
Ledger-IQ — finance close automation	Corporate Controller	Pilot	\$1.1M	\$0.9M (proj.)	On track
Compass — sales proposal drafting	CRO Office	Pilot	\$0.9M	\$1.4M (proj.)	Watch
Sentinel — fraud anomaly detection	Chief Risk Officer	Build	\$1.6M	\$2.2M (proj.)	On track
Archive-RAG — policy knowledge base	Chief Counsel	Scaled	\$0.7M	\$1.0M	On track
Foundry — agentic workflow platform	CTO Office	Build	\$2.1M	—	On track

Table 1. Tier-1 AI initiative portfolio as of 31 May 2026. Projected benefits are FP&A-validated; estimates; 'Watch' indicates a schedule or adoption risk under active management.

Outlook for the Second Half of 2026

The second half of the fiscal year is about depth, not breadth. We will hold the portfolio at roughly its current size while pushing the two scaled programs, Atlas and Beacon, toward full-population adoption, and graduating Ledger-IQ and Compass from pilot to scale subject to their stage-gate reviews in September. The Foundry agentic platform represents our most consequential technical bet: early internal trials suggest that supervised multi-step agents can compress procurement and vendor-onboarding cycle times by 40–60%, but they also raise the bar for evaluation rigor, rollback design, and audit logging. We will not scale agentic workflows until the Council's autonomy framework — currently in draft — is ratified.

Finally, talent remains the binding constraint. Attrition among senior machine-learning engineers is below industry benchmarks, but demand from business units now exceeds platform capacity by roughly three to one. The Q3 hiring plan adds eighteen platform roles, and we are piloting a rotational 'AI Guild' program that embeds certified practitioners from business units into platform squads for six-month tours. Leadership endorsement of that program at the July offsite is the single most valuable action this Council can take to sustain our momentum.